Experiment - 6

Student Name: Anshu Sharma   **UID:** 23BAI70204

Branch: BE-AIT-CSE   **Section/Group:** 23AIT_KRG-G1_A

Semester: 5th   **Date of Performance:** 24 Sept, 2025

Subject Name: ADBMS   **Subject Code:** 23CSP-333

1. Aim:

MEDIUM LEVEL PROBLEM:
To develop an automated database solution for TechSphere Solutions to monitor gender diversity. This will be achieved by creating a PostgreSQL stored procedure that takes a specific gender ('Male' or 'Female') as an input parameter, calculates the total count of employees for that gender, and returns the result as an output parameter for immediate HR reporting. This will allow the HR department to instantly and accurately retrieve the required diversity metrics.

HARD LEVEL PROBLEM:
The aim is to automate SmartShop's ordering and inventory management by creating a robust system that, upon a customer's request, verifies stock availability. If the stock is sufficient, the system will execute a single, atomic transaction to log the sale, update the product inventory, and confirm the transaction in real-time. Conversely, if the stock is insufficient, the system will reject the transaction and immediately inform the user, ensuring the integrity of both sales and inventory data.

2. Objective:

- Develop a PostgreSQL stored procedure to calculate and return the count of employees for a specified gender.
- Enable the HR department to instantly retrieve accurate gender diversity metrics for reporting.
- Automate the ordering process to verify stock availability for a customer's request.
- Execute a single, atomic transaction to log successful sales and update inventory levels in real-time.
- Implement a system to reject transactions with insufficient stock and provide immediate feedback to the user.
- Ensure the integrity of both sales and inventory data through a transactional approach.

3. Theory:

A **stored procedure** is a pre-compiled set of SQL statements and control flow logic stored in the database. Unlike a simple SQL query, a stored procedure can accept input parameters, perform conditional logic (like IF/ELSE), and return values. This makes them ideal for encapsulating complex, repeatable business logic. In the context of the TechSphere Solutions problem, a stored procedure provides an efficient and reusable way to solve a recurring HR task. Instead of writing a new SELECT and COUNT query every time, the HR team can simply call a pre-built procedure with the required gender as an argument. This approach improves data retrieval speed by reducing network traffic and, more importantly, ensures that the calculation is consistently and accurately performed every time, a crucial aspect of reliable reporting.

In database management, a **transaction** is a single unit of work that contains one or more operations. Transactions are essential for maintaining data integrity, especially in applications like inventory management where multiple related operations must succeed or fail together. This is governed by the ACID properties:

- **Atomicity:** This is the most critical concept for the SmartShop problem. Atomicity ensures that a transaction is treated as a single, indivisible unit. It either completes fully, with all its operations succeeding, or it fails completely, with none of its operations taking effect. This prevents partial updates that could lead to data inconsistencies—for example, a sale being logged without a corresponding reduction in inventory.
- **Consistency:** The database state remains valid before and after the transaction.
- **Isolation:** Concurrent transactions do not interfere with one another.
- **Durability:** Once a transaction is committed, its changes are permanent.

By using a transactional stored procedure, SmartShop can ensure that an order is only logged if the stock can be successfully updated. If the stock check fails, the transaction is rejected, and the database remains in its original, consistent state. This guarantees that the inventory count is always accurate and that no stock is "sold" if it's not physically available, which is vital for real-time business operations.

4. Procedure:

Medium Level Solution:

- **Setup:** Create an employee_info table and populate it with sample data, including employee names, genders, and other details.
- **Procedure Creation:** Develop a stored procedure named sp_get_employees_by_gender. This procedure takes a gender as an input parameter and an integer output parameter.
- **Business Logic:** Inside the procedure, a SELECT COUNT query counts all employees that match the input gender. The result is then stored in the output parameter.
- **Execution:** The procedure is called with a specific gender value (e.g., 'Male'), and a RAISE NOTICE command is used to print the final count, demonstrating a simple yet powerful automated reporting feature.

Hard Level Solution:

- **Setup:** Establish a database schema with products and sales tables to represent inventory and order history, respectively. Insert sample data into both tables.
- **Procedure Creation:** Create a stored procedure named pr_buy_products that accepts the product name and quantity as input.
- **Transactional Logic:** The procedure first checks if the requested quantity is available in the products table.
- **Conditional Processing:**
- **If sufficient stock:** The procedure executes a series of steps within a transaction: it inserts a new record into the sales table, updates the products table to reflect the reduced inventory (quantity_remaining) and increased sales (quantity_sold), and then prints a success message.
- **If insufficient stock:** The procedure immediately prints an "INSUFFICIENT QUANTITY" message without logging a sale or altering the inventory tables.
- **Execution:** Test the procedure with different values to demonstrate both a successful sale (when sufficient stock is available) and a failed transaction (when the quantity is too high), showcasing its transactional integrity and errorhandling capabilities.

5. Code:

-----------------------EXPERIMENT 06 (MEDIUM LEVEL)-----------------------

```sql
CREATE TABLE employee_info (
id SERIAL PRIMARY KEY, name
VARCHAR(50) NOT NULL, gender
VARCHAR(10) NOT NULL, salary
NUMERIC(10,2) NOT NULL, city
VARCHAR(50) NOT NULL
);

INSERT INTO employee_info (name, gender, salary, city)
VALUES
('Alok', 'Male', 50000.00, 'Delhi'),
('Priya', 'Male', 60000.00, 'Mumbai'),
('Rajesh', 'Female', 45000.00, 'Bangalore'),
('Sneha', 'Male', 55000.00, 'Chennai'),
('Anil', 'Male', 52000.00, 'Hyderabad'),
('Sunita', 'Female', 48000.00, 'Kolkata'),
('Vijay', 'Male', 47000.00, 'Pune'),
('Ritu', 'Male', 62000.00, 'Ahmedabad'),
('Amit', 'Female', 51000.00, 'Jaipur');




CREATE OR REPLACE PROCEDURE sp_get_employees_by_gender(
IN p_gender VARCHAR(50),
OUT p_employee_count INT
)
LANGUAGE plpgsql
AS $$
BEGIN

SELECT COUNT(id)
INTO p_employee_count
FROM employee_info
WHERE gender = p_gender;


RAISE NOTICE 'Total employees with gender %: %', p_gender, p_employee_count;
END;
$$;


CALL sp_get_employees_by_gender('Male', NULL);
```

------------------------EXPERIMENT 06 (HARD LEVEL)-------------------------

```sql
CREATE TABLE products (
product_code VARCHAR(10) PRIMARY KEY,
product_name VARCHAR(100) NOT NULL, price
NUMERIC(10,2) NOT NULL,
quantity_remaining INT NOT NULL,
quantity_sold INT DEFAULT 0
);

CREATE TABLE sales ( order_id
SERIAL PRIMARY KEY, order_date
```

```sql
DATE NOT NULL, product_code
VARCHAR(10) NOT NULL,
quantity_ordered INT NOT NULL,
sale_price NUMERIC(10,2) NOT NULL,
FOREIGN KEY (product_code) REFERENCES products(product_code)
);

INSERT INTO products (product_code, product_name, price, quantity_remaining,
quantity_sold)
VALUES
('P001', 'iPHONE 13 PRO MAX', 109999.00, 10, 0),
('P002', 'Samsung Galaxy S23 Ultra', 99999.00, 8, 0),
('P003', 'iPAD AIR', 55999.00, 5, 0),
('P004', 'MacBook Pro 14"', 189999.00, 3, 0),
('P005', 'Sony WH-1000XM5 Headphones', 29999.00, 15, 0);

INSERT INTO sales (order_date, product_code, quantity_ordered, sale_price)
VALUES
('2025-09-15', 'P001', 1, 109999.00),
('2025-09-16', 'P002', 2, 199998.00),
('2025-09-17', 'P003', 1, 55999.00),
('2025-09-18', 'P005', 2, 59998.00),
('2025-09-19', 'P001', 1, 109999.00);

SELECT * FROM PRODUCTS;
SELECT * FROM SALES;


CREATE OR REPLACE PROCEDURE pr_buy_products(
IN p_product_name VARCHAR,
IN p_quantity INT
)
LANGUAGE plpgsql
AS $$ DECLARE
v_product_code VARCHAR(20);
v_price FLOAT; v_count INT;
BEGIN


SELECT COUNT(*)
INTO v_count
FROM products
WHERE product_name = p_product_name
AND quantity_remaining >= p_quantity;


IF v_count > 0 THEN


SELECT product_code, price
INTO v_product_code, v_price
FROM products
WHERE product_name = p_product_name;


INSERT INTO sales (order_date, product_code, quantity_ordered, sale_price)
```

```
        VALUES (CURRENT_DATE, v_product_code, p_quantity, (v_price * p_quantity));
        UPDATE products
        SET quantity_remaining = quantity_remaining – p_quantity, quantity_sold
        = quantity_sold + p_quantity
        WHERE product_code = v_product_code;


        RAISE NOTICE 'PRODUCT SOLD..! Order placed successfully for % unit(s) of %.',
        p_quantity, p_product_name;

        ELSE

        RAISE NOTICE 'INSUFFICIENT QUANTITY..! Order cannot be processed for % unit(s)
        of %.', p_quantity, p_product_name;
        END IF;
        END;
        $$;


        CALL pr_buy_products ('MacBook Pro 14"', 1);
```
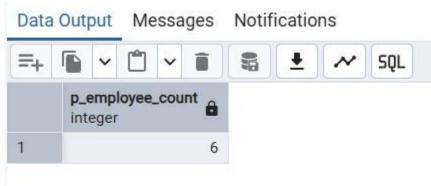
6. Output:

7. Learning Outcomes:

- **Encapsulation of Business Logic:** Learners will understand how stored procedures can be used to encapsulate complex business logic (e.g., counting employees by a specific criteria) into a single, reusable function. This promotes modularity and maintainability in database design.

- **Procedural Language Fundamentals:** This experiment provides a practical introduction to procedural languages like PL/pgSQL, including the use of DECLARE, BEGIN...END, input and output parameters, and the RAISE NOTICE command for debugging and reporting.

- **Efficient Reporting:** Learners will grasp how stored procedures can create a more efficient reporting layer. Instead of requiring a full query with aggregations from the client, the client only needs to make a simple CALL to the procedure, which handles the computation on the database server, reducing network overhead.

- **Encapsulation of Business Logic:** Learners will understand how stored procedures can be used to encapsulate complex business logic (e.g., counting employees by a specific criteria) into a single, reusable function. This promotes modularity and maintainability in database design.

- **Transactional Integrity (ACID Properties):** The primary learning outcome is a deep understanding of atomicity and transactional integrity. Learners will see firsthand how a single, cohesive transaction prevents data inconsistencies that could arise from partial updates, such as logging a sale without updating the inventory.

- **Real-time Inventory Management:** This experiment demonstrates a core function of ecommerce systems. Learners will understand how to use stored procedures to handle real-time inventory checks and updates, ensuring that stock levels are always accurate and preventing overselling.

- **Error Handling and User Feedback:** The use of IF/ELSE logic and RAISE NOTICE provides a clear example of how to handle different outcomes (success vs. failure) within a stored procedure and provide direct, real-time feedback to the user, a critical component of any user-facing application.